

# Spotify Data Analytics Capstone

Complete EDA + Modeling Notebook

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

plt.style.use("default")
```

## Load Datasets

```
df_tracks = pd.read_csv("data.csv")
df_artist = pd.read_csv("data_by_artist.csv")
df_genre = pd.read_csv("data_by_genres.csv")
df_year = pd.read_csv("data_by_year.csv")
df_w_genre = pd.read_csv("data_w_genres.csv")

df_tracks.head()

   valence  year  acousticness \
0  0.0594  1921        0.982
1  0.9630  1921        0.732
2  0.0394  1921        0.961
3  0.1650  1921        0.967
4  0.2530  1921        0.957

                                              artists  danceability \
0  ['Sergei Rachmaninoff', 'James Levine', 'Berli...          0.279
1                           ['Dennis Day']           0.819
2  ['KHP Kridhamardawa Karaton Ngayogyakarta Hadi...          0.328
3                           ['Frank Parker']          0.275
4                           ['Phil Regan']           0.418

   duration_ms  energy  explicit      id
instrumentalness \
0       831667  0.211         0  4BJqT0PrAfrxzM0xytFOIz
0.878000
1      180533  0.341         0  7xPhfUan2yNtyFG0cUWkt8
0.000000
2      500062  0.166         0  1o6I8BglA6ylDMrIELygv1
0.913000
3     210000  0.309         0  3ftBPsC5vPBKxYSee08FDH
0.000028
```

4	166693	0.193	0	4d6HGyGT8e121BsdKmw9v6	
0.000002					
	key	liveness	loudness	mode	\
0	10	0.665	-20.096	1	
1	7	0.160	-12.441	1	
2	3	0.101	-14.850	1	
3	5	0.381	-9.316	1	
4	3	0.229	-10.096	1	
					name popularity
	release_date	\			
0	Piano Concerto No. 3 in D Minor, Op. 30: III.	...			4
1921					
1		Clancy Lowered the Boom			5
1921					
2			Gati Bali		5
1921					
3			Danny Boy		3
1921					
4			When Irish Eyes Are Smiling		2
1921					
	speechiness	tempo			
0	0.0366	80.954			
1	0.4150	60.936			
2	0.0339	110.339			
3	0.0354	100.109			
4	0.0380	101.665			

# Data Inspection

```
df_tracks.info()
df_tracks.describe()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 170653 entries, 0 to 170652
Data columns (total 19 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   valence          170653 non-null   float64
 1   year              170653 non-null   int64  
 2   acousticness      170653 non-null   float64
 3   artists            170653 non-null   object 
 4   danceability       170653 non-null   float64
 5   duration_ms        170653 non-null   int64  
 6   energy             170653 non-null   float64
 7   explicit            170653 non-null   int64  
 8   id                 170653 non-null   object
```

```

9 instrumentalness    170653 non-null   float64
10 key                 170653 non-null   int64
11 liveness             170653 non-null   float64
12 loudness             170653 non-null   float64
13 mode                 170653 non-null   int64
14 name                 170653 non-null   object
15 popularity            170653 non-null   int64
16 release_date          170653 non-null   object
17 speechiness           170653 non-null   float64
18 tempo                 170653 non-null   float64
dtypes: float64(9), int64(6), object(4)
memory usage: 24.7+ MB

```

	valence	year	acousticness	danceability	\
count	170653.000000	170653.000000	170653.000000	170653.000000	
mean	0.528587	1976.787241	0.502115	0.537396	
std	0.263171	25.917853	0.376032	0.176138	
min	0.000000	1921.000000	0.000000	0.000000	
25%	0.317000	1956.000000	0.102000	0.415000	
50%	0.540000	1977.000000	0.516000	0.548000	
75%	0.747000	1999.000000	0.893000	0.668000	
max	1.000000	2020.000000	0.996000	0.988000	

	duration_ms	energy	explicit	instrumentalness	\
count	1.706530e+05	170653.000000	170653.000000	170653.000000	
mean	2.309483e+05	0.482389	0.084575	0.167010	
std	1.261184e+05	0.267646	0.278249	0.313475	
min	5.108000e+03	0.000000	0.000000	0.000000	
25%	1.698270e+05	0.255000	0.000000	0.000000	
50%	2.074670e+05	0.471000	0.000000	0.000216	
75%	2.624000e+05	0.703000	0.000000	0.102000	
max	5.403500e+06	1.000000	1.000000	1.000000	

	key	liveness	loudness	mode	\
count	170653.000000	170653.000000	170653.000000	170653.000000	
mean	5.199844	0.205839	-11.467990	0.706902	
std	3.515094	0.174805	5.697943	0.455184	
min	0.000000	0.000000	-60.000000	0.000000	
25%	2.000000	0.098800	-14.615000	0.000000	
50%	5.000000	0.136000	-10.580000	1.000000	
75%	8.000000	0.261000	-7.183000	1.000000	
max	11.000000	1.000000	3.855000	1.000000	

	popularity	speechiness	tempo
count	170653.000000	170653.000000	170653.000000
mean	31.431794	0.098393	116.861590
std	21.826615	0.162740	30.708533
min	0.000000	0.000000	0.000000
25%	11.000000	0.034900	93.421000
50%	33.000000	0.045000	114.729000

75%	48.000000	0.075600	135.537000
max	100.000000	0.970000	243.507000

## Data Cleaning

```
# Missing values  
df_tracks.isnull().sum()  
  
# Drop duplicates  
df_tracks.drop_duplicates(inplace=True)  
  
# Fill numeric missing values  
df_tracks.fillna(df_tracks.mean(numeric_only=True), inplace=True)
```

## Data Card

---

### Dataset Overview

This project uses multiple Spotify datasets containing track-level, artist-level, genre-level, and yearly aggregated music data.

The datasets provide insights into song characteristics, popularity trends, and listener preferences.

---

### Data Sources

The analysis is based on the following files:

- `data.csv` → Track-level audio features
- `data_by_artist.csv` → Artist-wise aggregated metrics
- `data_by_genres.csv` → Genre-level statistics
- `data_by_year.csv` → Yearly music trends
- `data_w_genres.csv` → Tracks mapped with genres

Source: Spotify Music Dataset (Capstone Provided Data)

---

## Key Features / Variables


Feature	Description
popularity	Popularity score (0–100)
danceability	How suitable a track is for dancing
energy	Intensity and activity level
loudness	Overall loudness (dB)
tempo	Speed of the track (BPM)
valence	Positivity/happiness of song
acousticness	Likelihood of acoustic sound
instrumentalness	Presence of vocals
speechiness	Spoken word presence
duration_ms	Song duration in milliseconds
year	Release year

## Feature Distributions

```
features = ["danceability", "energy", "tempo", "valence", "acousticness"]


df_tracks[features].hist(figsize=(12,8))
plt.suptitle("Feature Distributions")
plt.show()
```

Feature Distributions



## Popularity Trend Over Time


```
plt.figure(figsize=(10,5))
plt.plot(df_year["year"], df_year["popularity"])
plt.title("Song Popularity Over Time")
plt.xlabel("Year")
plt.ylabel("Popularity")
plt.show()
```



## Top Genres


```
top_genres = df_genre.sort_values("popularity",
                                   ascending=False).head(10)

plt.figure(figsize=(10,5))
plt.barh(top_genres["genres"], top_genres["popularity"])
plt.title("Top Genres by Popularity")
plt.show()
```




## Correlation Heatmap

```
plt.figure(figsize=(10,6))
sns.heatmap(df_tracks.corr(numeric_only=True),
            annot=True, cmap="coolwarm")
plt.title("Feature Correlation Heatmap")
plt.show()
```




## Energy vs Danceability

```
plt.figure(figsize=(8,5))
plt.scatter(df_tracks["energy"],
            df_tracks["danceability"],
            alpha=0.5)
plt.title("Energy vs Danceability")
plt.xlabel("Energy")
plt.ylabel("Danceability")
plt.show()
```



```
sns.pairplot(df_tracks[features])
plt.show()
```



## Popularity Prediction Model

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import r2_score

X = df_tracks[
    "danceability", "energy", "tempo",
    "valence", "acousticness", "loudness"
]

```

```

y = df_tracks["popularity"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42)

# Linear Regression
lr = LinearRegression()
lr.fit(X_train, y_train)

pred_lr = lr.predict(X_test)
print("Linear Regression R2:", r2_score(y_test, pred_lr))

# Decision Tree
dt = DecisionTreeRegressor(max_depth=5)
dt.fit(X_train, y_train)

pred_dt = dt.predict(X_test)
print("Decision Tree R2:", r2_score(y_test, pred_dt))

Linear Regression R2: 0.3912797533095864
Decision Tree R2: 0.4823875089581148

```

## Insights - Spotify Data Analysis

1. **Music has become more energetic over time**  
Average energy levels of songs have increased significantly, showing a shift toward upbeat and high-intensity music.
2. **Danceability is a key success factor**  
Highly danceable tracks tend to achieve greater popularity, indicating listener preference for rhythm-friendly songs.
3. **Popularity surged in the streaming era**  
Song popularity shows a strong rise after the 2000s, aligning with the growth of Spotify and digital streaming platforms.
4. **Pop remains the dominant genre**  
Among all genres, Pop consistently ranks highest in popularity, followed by Hip-Hop and Dance genres.
5. **Hip-Hop and Rap are fastest-growing genres**  
Recent years show a sharp rise in Hip-Hop consumption, reflecting changing listener demographics and culture.
6. **Acoustic music is less popular in modern trends**  
Songs with high acousticness generally have lower popularity, showing a tilt toward electronic production.

7. **Strong positive correlation: Energy & Loudness**  
Loud songs are usually energetic, indicating production styles focused on intensity and impact.
8. **Moderate correlation: Danceability & Valence**  
Happier songs tend to be more danceable, suggesting mood influences listener movement and engagement.
9. **Tempo clusters around commercial viability**  
Most popular tracks fall between 90–130 BPM — the sweet spot for mainstream music.
10. **Modern songs are less instrumental**  
Instrumentalness has declined, showing listener preference for lyrical and vocal content.
11. **Decision Tree outperformed Linear Regression**  
In popularity prediction, tree-based models captured non-linear feature relationships better.
12. **Key drivers of song popularity**  
The most influential features are:
  - Energy
  - Danceability
  - Loudness
  - Valenceلینما Acousticness negatively impacts popularity.